

Machine Learning


“Field of study that gives computers the ability to learn without being explicitly programmed” - Arthur Samuel

1. Supervised learning


Let's start from an example with a dataset

Living Area (sq ft)	Price (\$)
2104	400
1600	330
2400	232


We can represent the data as follows



The common approach to this kind of problems is finding a *function* that operates with the data



In machine learning, the mapping is learnt automatically, and this leads to a different approach that is *data dependent*



If the output data is real valued, we are considering a *regression*, if it's binary we call it *classification*. This distinction is implicated in the trained data.

Example

- Linear regression: regression
- Logistic regression: classification

2. Linear regression

	Living area	# Bedrooms	Price
i			

x
y

N: number of examples (rows)

D: numbers of features (records)

\vec{x}^i : i^{th} input $\in \mathbb{R}^d$

y^i : i^{th} input $\in \mathbb{R}$

Let's define a function H related to the *ordinary least squares assumption*

$$H_{\theta}(\vec{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

θ : parameters, $\theta \in \mathbb{R}$

$\theta_1, \dots, \theta_d$: weights

θ_0 : bias

In general

$$H_{\theta}(x) = \sum_{i=0}^d \theta_i x_i$$

and the input of this function is one row at a time.

Our mission is to determine θ with the learning algorithm. To do this, we use a *cost function* J that associates greater thetas to a minor cost and vice versa.

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n [H_{\theta}(x^i) - y^i]^2 \in \mathbb{R}$$

What is happening here is that we are *comparing* our cost function with the right answer from y vector.

Now what we want is to *optimize* θ and, in linear regressions, there are two main algorithms to do that:

- LMS algorithm
- Normal equation

I. LMS algorithm


It is an iterative algorithm. Let's set an initial value for θ (any).

$$\theta := 0$$

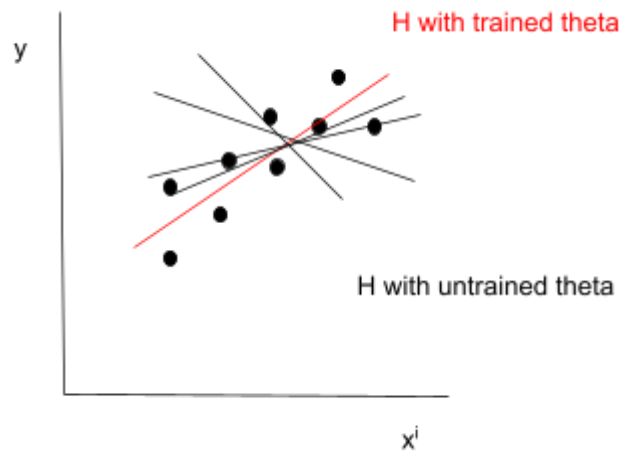
We need to repeat the following operation until *convergence* (not necessarily happening), for every element of theta (vector) and obtain a *new corresponding* θ .

$$\left\{ \theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \right\}$$

With θ_j being the j^{th} of the *gradient*



The sign is flipped, α is the *learning grade* and scales the vector (a fixed value in this case), step by step it will eventually converge. The circles (ellipses) represent the set of points of θ having the same cost. Now let's take the new θ and use it in our function H .



II. The normal equation

Uses the dataset from the beginning, instead of iterating gradually.

$$X = \begin{bmatrix} & \\ & \\ & \end{bmatrix} : (d+1 \times n)\text{-matrix}$$

$$\vec{y} = \begin{bmatrix} \\ \\ \end{bmatrix} : n\text{-vector}$$

$$\theta = \begin{bmatrix} \\ \\ \end{bmatrix} : d+1\text{-vector}$$

$$X\theta - \vec{y} = \begin{bmatrix} \theta^{x^i} - y^i \\ \dots \end{bmatrix}$$

Starting from the definition of J , we can work our way out to another model

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n [H_{\theta}(x^i) - y^i]^2 \in \mathbb{R}$$

$$J(\theta) = \frac{1}{2} (X\theta - \vec{y})^T (X\theta - \vec{y})$$

$$\nabla_{\theta} J(\theta) = 0$$

$$\nabla_{\theta} \frac{1}{2} (X\theta - \vec{y})^T (X\theta - \vec{y}) = 0$$


$$\nabla_{\theta} \frac{1}{2} [2X^T X\theta - 2X^T \vec{y}] = 0$$

$$X^T X\theta - X^T \vec{y} = 0$$

We come to the final conclusion, which is the solution for θ for a few “lucky” ML models.

$$\theta = (X^T X)^{-1} X^T \vec{y}$$

Now, let's take a statistical approach, and analyze the *noise* related to this kind of regression.



We can define y^i as follows

$$y^i = \theta^T(x^i) + \varepsilon^i$$

Where $\theta^T(x^i)$ is the *HP function* and ε^i is the *noise randomness* and $\varepsilon^i \sim N(0, \sigma^i)$.

Let's better define epsilon

$$\varepsilon^i = \bar{y}^i - \varepsilon^i(x^i)$$

And the associated *gaussian density distribution function* P

$$P(\varepsilon^i) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ -\frac{1}{2} \frac{(\varepsilon^i - d^2)}{\sigma^2} \right\}$$

$$P(\bar{y}^i || x^i i \theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ \frac{-[y^i - \theta^T(x^i)]^2}{2\sigma^2} \right\}$$

$$P(y || x i \theta) = \prod_{i=1}^n P(y^i || x^i i \theta)$$

Let's define L function of likelihood, in order to *maximize* the result

$$L(\theta) = \prod_{i=1}^n P(y^i || x^i i \theta)$$

$$\log L(\theta) = \log \prod_{i=1}^n P(y^i || x^i i \theta) = \sum_{i=1}^n \log P(y^i || x^i i \theta)$$

$$\sum_{i=1}^n \log \left[\frac{1}{\sqrt{2\pi}\sigma} \exp \left\{ \frac{-[y^i - \theta^T(x^i)]^2}{2\sigma^2} \right\} \right]$$

$$\sum_{i=1}^n \left(\log \frac{1}{\sqrt{2\pi}\sigma} \right) + \left\{ \frac{-(y^i - \theta(x^i))^2}{2\sigma^2} \right\}$$

$$\sum_{i=1}^n \left\{ K + \frac{-(y^i - \theta(x^i))^2}{2\sigma^2} \right\}$$

Coming to the conclusion

$$\log L(\theta) = nK - \frac{1}{2\sigma^2} \sum_{i=1}^n (y^i - \theta^T(x^i))^2$$