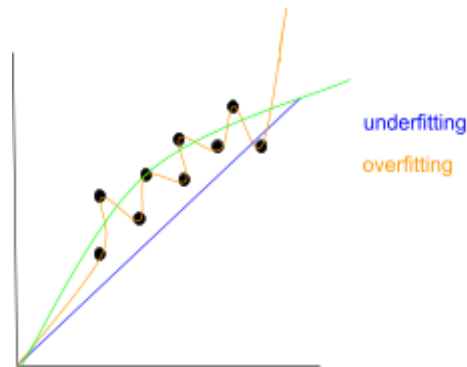


Machine Learning

“Field of study that gives computers the ability to learn without being explicitly programmed” - Arthur Samuel

Example “fitting” over a linear regression



- underfitting

$$y = \theta_0 + \theta_1 x_1$$

- overfitting

$$y = \sum_{i=0}^d \theta_j x_1^j$$

- best size

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_1^2$$

1. Locally weighted regression

Model	Loss function	Output
Linear	$\sum_i (y^{(i)} - \theta^T x^{(i)})^2$	$\theta^T x$
Locally weighted (LWR)	$\sum_i w^{(i)} (y^{(i)} - \theta^T x^{(i)})$	$\theta^T x$

Where

$$w^{(i)} = \exp \left(\frac{-(x^{(i)} - x)^2}{2\tau^2} \right)$$

And $x^{(i)}$ being a *train example*, x a *test example*, τ the *bandwidth parameter*.

Note that $0 \leq w^{(i)} \leq 1$ and for examples very far from each other it tends to 0, for examples very close to each other it tends to 1

In LWR, the output $\theta^T x$ gets re-estimated for every test example. You calculate a new weight for the specific test example (each one), so the bigger the data set, the bigger the value of the model.

2. Logistic regression (classification)

When we talked about the linear regression, we meant a model such that

$$\mathbb{R}^n \rightarrow \mathbb{R}$$

While in linear classification, we are considering a *restricted* set of $y^{(i)}$ s

$$\mathbb{R}^n \rightarrow \{0, 1\}$$

The previous example refers explicitly to a binary classification, but it could be a different set.

Linear regression	Logistic regression
$y \in \mathbb{R}$	$y \in \{0, 1\}$
$h_{\theta}(x) = \theta^T x$	$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + \exp(-\theta^T x)}$


The classification's hypothesis function is called *logistic function*

$$h_{\theta}^{logistic} = \frac{1}{1 + \exp(-\theta^T x)}$$

So the actual score $\theta^T x$ determines if the output of the function and

$$0 \leq g(\theta^T x) \leq 1$$

Ideally, this mean that the function *shrinks* all the possible outcomes in a $\{0, 1\}$ range.



Now, let's talk about probabilities

$$\begin{cases} P(y = 1 | x_1 \theta) = h_\theta(x) \\ P(y = 0 | x_1 \theta) = 1 - h_\theta(x) \end{cases}$$

is a *Bernoulli distribution*.

x could be any transformation (x^2, x^n, \dots)

Let's rearrange the terms and we obtain

$$P(y | x_1 \theta) = h_\theta(x)^y (1 - h_\theta(x))^{1-y}$$

the probability distribution function.

Let's now define $L(\theta)$ the likelihood function

$$L(\theta) = P(y | x_1 \theta) = \prod_{i=1}^n P(y^{(i)} | x_1^{(i)} \theta)$$

$$L(\theta) = \prod_{i=1}^n h_\theta(x^{(i)})^{y^{(i)}} [1 - h_\theta(x^{(i)})]^{1-y^{(i)}}$$

$$\log L(\theta) = \sum_{i=1}^n \left\{ y^{(i)} \log h_\theta(x^{(i)}) + (1 - y^{(i)}) \log [1 - h_\theta(x^{(i)})] \right\}$$

Now, let's apply *gradient ascent* to $L(\theta)$ in order to maximize the likelihood

$$\begin{cases} \theta = 0 \\ \theta_k = \theta_{k-1} + \alpha \nabla_\theta \log L(\theta) \end{cases} \quad \text{initialization}$$

$$\nabla_\theta \log L(\theta) = \sum_{i=1}^n y^{(i)} \nabla_\theta \log \frac{1}{1 + \exp(-\theta^T x)} + (1 - y^{(i)}) \nabla_\theta \log \left(1 - \frac{1}{1 + \exp(-\theta^T x)} \right)$$

The gradient of the function is

$$\nabla_\theta L(\theta) = \sum_{i=1}^n (y^{(i)} - h_\theta(x^{(i)})) x^{(i)}$$

Now, repeat the algorithm until convergence and make predictions by calculating $\theta^T x$ (we want it to be \gg or \ll , so the prediction is either 1 or 0).

3. Newton's method

NsM is an optimization algorithm, very efficient to find zeros of a function. In respect to gradient descent, Newton's converges faster if the considered starting point is close to the solution. In this case, we say Newton's is more efficient since the 0 of a gradient might or not be a *minimum*.

In general, the Newton's algorithm looks like this

$$\theta_k = \theta_{k-1} - \frac{f(\theta)}{f'(\theta)}$$

And iterate until convergence $f(\theta) = 0$

Now let's apply the gradient, set $f = L'$ and we obtain

$$\theta_k = \theta_{k-1} - \frac{L'(\theta)}{L''(\theta)}$$

Now, let's consider an \mathbb{R}^n vector space and we finally obtain the expression to iterate

$$\theta_k = \theta_{k-1} - \alpha H^{-1} \nabla_{\theta} L(\theta)$$

where H^{-1} is the *hessian inverse*.

$$H_{ij} = \frac{\partial^2 L(\theta)}{\partial \theta_i \partial \theta_j}$$

Newton's always puts you into the closest *max* or *min* in the vector space. This method is really effective for small values of n , while it is not for higher values.

4. Stochastic gradient descent

In *linear regression* we have defined the *train loss function* $TrainLoss(\theta)$ as follows

$$TrainLoss = \sum_{i=0}^n J(\theta, x^{(i)}, y^{(i)})$$

where

$$J(\theta, x^{(i)}, y^{(i)}) = (\theta^T x^{(i)} - y^{(i)})^2$$

So the gradient descent's update rule would look like this

$$\theta_k = \theta_{k-1} - \alpha \sum_{i=0}^n \nabla_{\theta} J(\theta, x^{(i)}, y^{(i)})$$

Which can be computationally expensive and inefficient if the data set is big.

The *SGD* update rule, instead of taking the sum of the gradients of all the train examples with respect to θ and iterating this process until convergence, uses the following function

$$\tilde{L}(\theta) = J(\theta, x^{(r)}, y^{(r)}), \quad r \sim [0, 1]$$

that takes one *random* example at a time.

This is much more efficient for big data sets. It is also important to mention that the idea of *convergence* is also different in SGD. Here the convergence means that the gradients reached a point that is *within a circle around the converge point* and the radius of the circle equals to *learning grade* α (or step size).

